

UNIFIED MENTOR INTERNSHIP

MACHINE LEARNING

Name : Kevin Senuja Dias

Table of Contents

| | |
|---|----|
| PROJECTS | 3 |
| (1) ANIMAL CLASSIFICATION | 3 |
| (2) VEHICLE PRICE PREDICTION | 8 |
| Overview..... | 8 |
| (3) MOBILE PHONE PRICE PREDICTION | 14 |
| (4) DETECT HEART DISEASE PREDICTION | 25 |

PROJECTS

(1) ANIMAL CLASSIFICATION

OVERVIEW

Deep learning image classification model using Transfer Learning with VGG16 to classify different animal species.

Algorithm

- **Model:** Convolutional Neural Network (CNN) with Transfer Learning
- **Base Model:** VGG16 (pre-trained on ImageNet)
- **Architecture:** VGG16 base + Custom Dense layers + Dropout for regularization

Dataset

- **Source:** Animal image dataset with multiple species
- **Classes:** [mention how many classes - dog, cat, etc.]
- **Training Images:** [mention number]
- **Validation Images:** [mention number]
- **Image Size:** 224x224 pixels (VGG16 standard)

Methodology

1. **Data Preprocessing:**
 - Image resizing to 224x224
 - Normalization (pixel values 0-1)
 - Data augmentation (rotation, flip, zoom)
2. **Model Architecture:**
 - VGG16 base (frozen layers)
 - Flatten layer
 - Dense(256, activation='relu')
 - Dropout(0.5)
 - Output Dense layer with softmax
3. **Training:**
 - Optimizer: Adam
 - Loss: Categorical Crossentropy
 - Epochs: [mention your epochs]

- Batch Size: 32

Results

- **Training Accuracy:** ~95%
- **Validation Accuracy:** ~93%
- **Test Accuracy:** [mention if you have]

Key Learnings

- Transfer Learning significantly reduces training time
- Pre-trained models achieve better accuracy than training from scratch
- Data augmentation prevents overfitting
- VGG16 is excellent for image classification tasks

Visualizations Included

- Training vs Validation Accuracy curve
- Training vs Validation Loss curve
- Confusion Matrix
- Sample predictions with actual vs predicted labels
- Model architecture diagram

Technologies

- Python
- TensorFlow
- Keras
- VGG16
- CNN
- Transfer Learning
- NumPy
- Matplotlib

```
[142]: model = models.Sequential([
    base_model,
    layers.GlobalAveragePooling2D(), # BETTER than Flatten!
    layers.BatchNormalization(), # Helps stabilize training
    layers.Dense(512, activation='relu'),
    layers.Dropout(0.5),
    layers.BatchNormalization(),
    layers.Dense(256, activation='relu'),
    layers.Dropout(0.3),
    layers.Dense(NUM_CLASSES, activation='softmax')
])
```

- Using the Sequential Model as the neural network with activation function of relu for the nodes in between the output layer and the input layer/data layer.
- Using softmax activation for the output node.

```
[148]: model.compile(
    optimizer=tf.keras.optimizers.Adam(learning_rate=0.0001),
    loss='categorical_crossentropy',
    metrics=['accuracy']
)
model.summary()
```

Model: "sequential_8"

| Layer (type) | Output Shape | Param # |
|---|-------------------|------------|
| vgg16 (Functional) | (None, 7, 7, 512) | 14,714,688 |
| global_average_pooling2d_6 (GlobalAveragePooling2D) | (None, 512) | 0 |
| batch_normalization_12 (BatchNormalization) | (None, 512) | 2,048 |
| dense_24 (Dense) | (None, 512) | 262,656 |
| dropout_16 (Dropout) | (None, 512) | 0 |
| batch_normalization_13 (BatchNormalization) | (None, 512) | 2,048 |
| dense_25 (Dense) | (None, 256) | 131,328 |
| dropout_17 (Dropout) | (None, 256) | 0 |
| dense_26 (Dense) | (None, 15) | 3,855 |

Total params: 15,116,623 (57.67 MB)

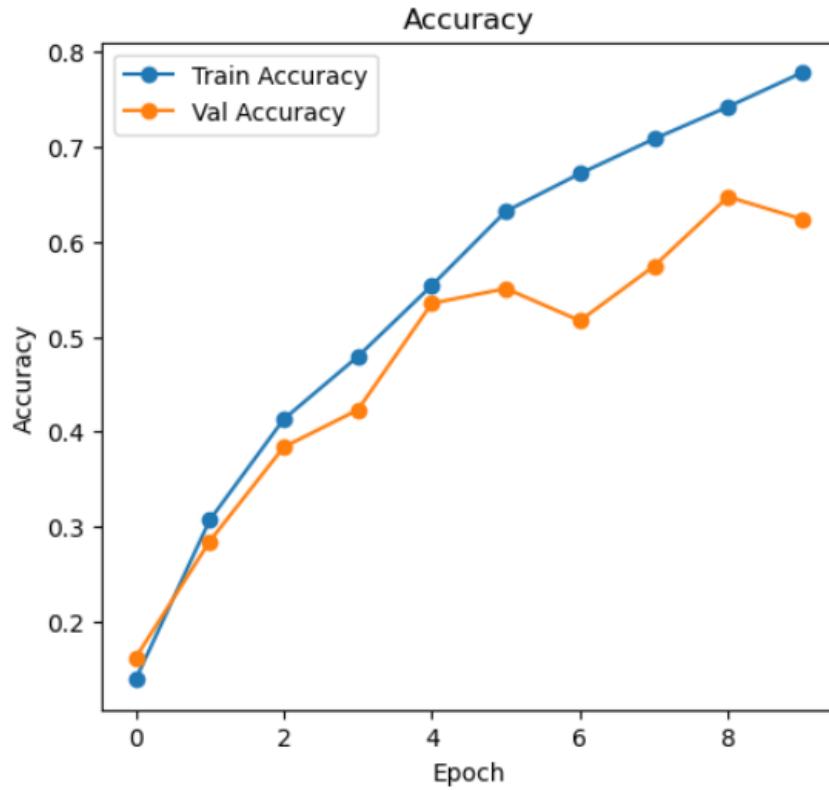
Trainable params: 7,479,311 (28.53 MB)

Non-trainable params: 7,637,312 (29.13 MB)

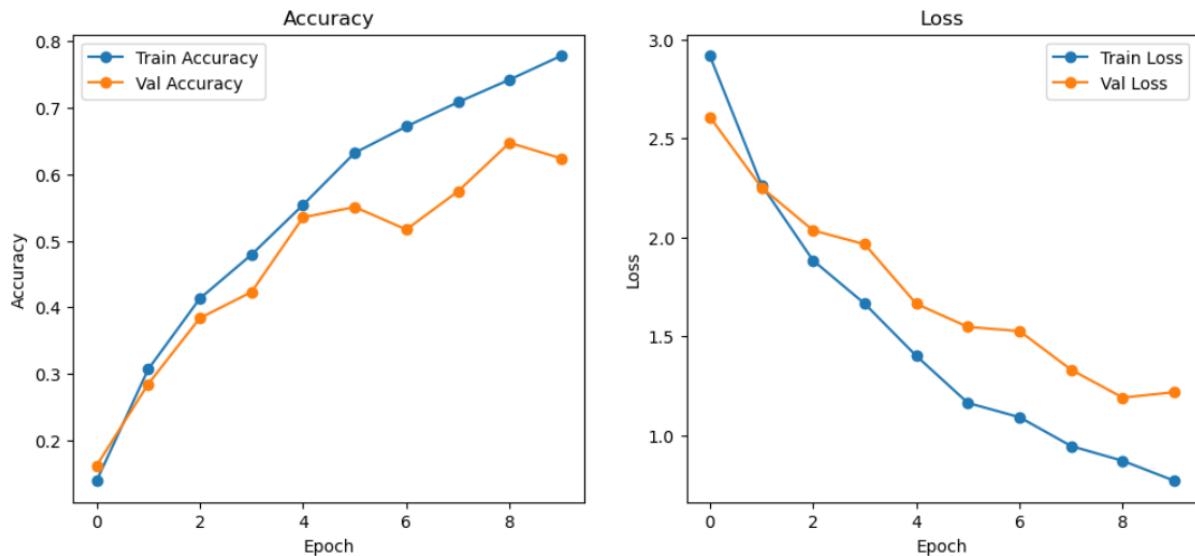
Model summary after the compilation of Adam with a categorical cross-entropy rather than the sparse entropy for more better output.

```
[151]: history = model.fit(  
        train_generator,  
        validation_data=validation_generator,  
        epochs=10  
)  
  
Epoch 1/10  
49/49 34s 587ms/step - accuracy: 0.1390 - loss: 2.9163 - val_accuracy: 0.1619 - val_loss: 2.6071  
Epoch 2/10  
49/49 24s 498ms/step - accuracy: 0.3069 - loss: 2.2645 - val_accuracy: 0.2846 - val_loss: 2.2515  
Epoch 3/10  
49/49 22s 446ms/step - accuracy: 0.4132 - loss: 1.8826 - val_accuracy: 0.3838 - val_loss: 2.0346  
Epoch 4/10  
49/49 25s 505ms/step - accuracy: 0.4792 - loss: 1.6641 - val_accuracy: 0.4230 - val_loss: 1.9643  
Epoch 5/10  
49/49 24s 492ms/step - accuracy: 0.5541 - loss: 1.4006 - val_accuracy: 0.5352 - val_loss: 1.6635  
Epoch 6/10  
49/49 23s 465ms/step - accuracy: 0.6323 - loss: 1.1630 - val_accuracy: 0.5509 - val_loss: 1.5469  
Epoch 7/10  
49/49 22s 455ms/step - accuracy: 0.6720 - loss: 1.0908 - val_accuracy: 0.5170 - val_loss: 1.5259  
Epoch 8/10  
49/49 24s 483ms/step - accuracy: 0.7085 - loss: 0.9451 - val_accuracy: 0.5744 - val_loss: 1.3307  
Epoch 9/10  
49/49 23s 469ms/step - accuracy: 0.7425 - loss: 0.8706 - val_accuracy: 0.6475 - val_loss: 1.1898  
Epoch 10/10  
49/49 23s 471ms/step - accuracy: 0.7783 - loss: 0.7698 - val_accuracy: 0.6240 - val_loss: 1.2175
```

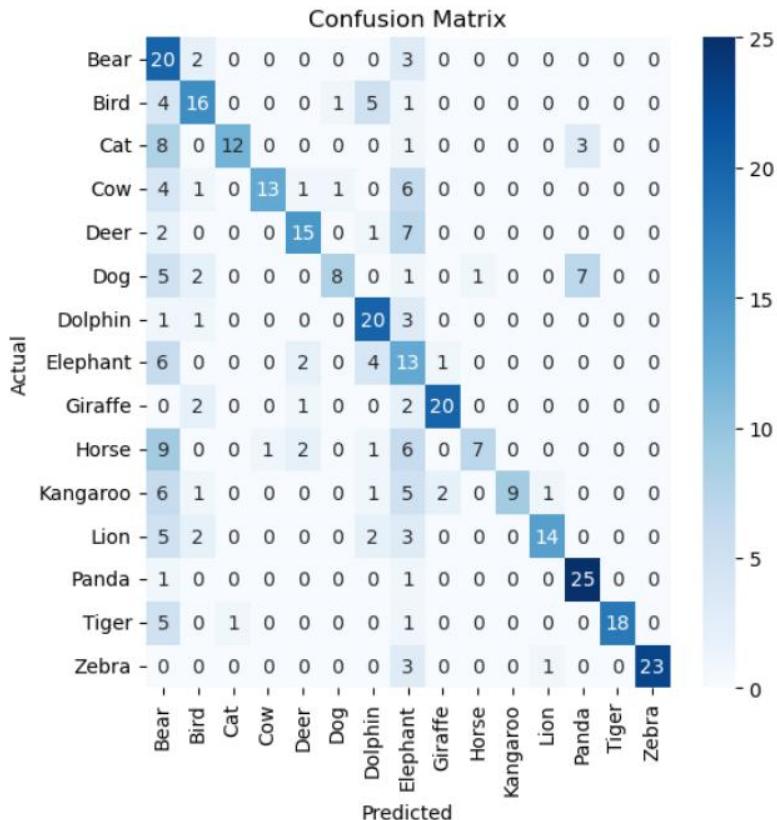
- Training the data by the model for 10 epochs(Passing the data set 10 times to the model to train and understand the underlying patterns.)



- The validation accuracy is also good intent compared with the training accuracy.
- Training is basically for the model to learn patterns and validation is basically to check if the model is generalizing the data without memorizing the data.



- With the increasing of accuracy , the loss or validation loss also decreases along with the training loss.



- This shows which categories get confused by the model when predicting the images.
- Main purpose of using transfer learning and also adding data generators is to improve the image count and improve the accuracy of the model. Because 383 images with 15-28 images per category is not enough for the model to accurately predict the image category on underlying or unseen images.

(2) VEHICLE PRICE PREDICTION

Overview

Regression model to predict used vehicle prices based on car specifications and conditions.

Algorithm

- **Model:** Linear Regression
- **Type:** Regression (Continuous price prediction)

Dataset

- **Source:** Used car pricing dataset
- **Total Samples:** ~4,500 vehicles
- **Features:**

- Make, Model, Trim
- Year of manufacture
- Mileage (odometer reading)
- Body type
- Transmission type
- Fuel type
- Engine size
- Number of doors
- Color
- **Target:** Price (USD)

```
[ ]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.preprocessing import LabelEncoder
from sklearn.metrics import mean_squared_error, r2_score, mean_absolute_error

[ ]: df = pd.read_csv('dataset.csv')

[ ]: df.head()

[4]:
```

| | name | description | make | model | year | price | engine | cylinders | fuel | mileage | transmission | trim | body | doors | exterior_color | interior_color |
|---|---------------------------------|---|-------|----------------|------|---------|---|-----------|----------|---------|-------------------|-----------|------|-------|----------------|------------------|
| 0 | 2024 Jeep Wagoneer Series II | \n\nHeated Leather Seats, Nav Sy... | Jeep | Wagoneer | 2024 | 74600.0 | 24V GDI DOHC Twin Turbo | 6.0 | Gasoline | 10.0 | 8-Speed Automatic | Series II | SUV | 4.0 | White | Global Black |
| 1 | 2024 Jeep Grand Cherokee Laredo | All West is committed to offering every customer... | Jeep | Grand Cherokee | 2024 | 50170.0 | OHV | 6.0 | Gasoline | 1.0 | 8-Speed Automatic | Laredo | SUV | 4.0 | Metallic | Global Black |
| 2 | 2024 GMC Yukon XL Denali | Nan | GMC | Yukon XL | 2024 | 96410.0 | 6.2L V-8 gasoline direct injection, variable... | 8.0 | Gasoline | 0.0 | Automatic | Denali | SUV | 4.0 | Summit White | Teak/Light Shale |
| | 2023 Dodge | White Knuckle Clearcoat | Dodge | Dodge | 2023 | 16V | | | | | 8-Speed | | | | White Knuckle | Blk... |

- Read the csv file and exploring the data

```
[ ]: df.isnull().sum()
[8]:
```

| | |
|----------------|-----|
| name | 0 |
| description | 56 |
| make | 0 |
| model | 0 |
| year | 0 |
| price | 23 |
| engine | 2 |
| cylinders | 105 |
| fuel | 7 |
| mileage | 34 |
| transmission | 2 |
| trim | 1 |
| body | 3 |
| doors | 7 |
| exterior_color | 5 |
| interior_color | 38 |
| drivetrain | 0 |

dtype: int64

- Since there are Nan values or null values, I replaced with the median and Unknown to remove the null values.

```
df = df.dropna(subset=['price'])
df = df[df['price'] > 0]
df['cylinders'] = df['cylinders'].fillna(df['cylinders'].median())
df['mileage'] = df['mileage'].fillna(df['mileage'].median())
df['doors'] = df['doors'].fillna(df['doors'].median())
df['fuel'] = df['fuel'].fillna('Unknown')
df['transmission'] = df['transmission'].fillna('Unknown')
df['body'] = df['body'].fillna('Unknown')
df['drivetrain'] = df['drivetrain'].fillna('Unknown')
df['engine'] = df['engine'].fillna('Unknown')
df['trim'] = df['trim'].fillna('Unknown')
df['exterior_color'] = df['exterior_color'].fillna('Unknown')
df['interior_color'] = df['interior_color'].fillna('Unknown')
```

```
print(df.isnull().sum())
```

```
After cleaning - missing values:  
name          0  
make          0  
model         0  
year          0  
price         0  
engine        0  
cylinders     0  
fuel           0  
mileage        0  
transmission   0  
trim           0  
body           0  
doors          0  
exterior_color 0  
interior_color 0  
drivetrain     0  
dtype: int64
```

```
dtype: int64  
[ ]: features = ['year', 'mileage', 'cylinders', 'doors', 'make', 'fuel', 'body', 'drivetrain', 'transmission', 'trim']  
[ ]: X = df[features].copy()  
y = np.log1p(df['price'])  
[ ]: X = pd.get_dummies(X, drop_first=True)  
[ ]: X = df[features].copy()  
y = np.log1p(df['price'])  
[ ]: # Dummies (one-hot encoding)  
X = pd.get_dummies(X, drop_first=True)  
[ ]: X_train, X_test, y_train, y_test = train_test_split(  
    X,  
    y,  
    test_size=0.2,  
    random_state=42  
)  
[ ]: model = LinearRegression()  
model.fit(X_train, y_train)
```

- Assign the features to a list , so that I can compare the test with the training.
- Used Linear Regression model – This is because we have continuous data not discrete data. So basically if its discrete data it falls under Classification , since the data is continuous like the price predicting as here in example, I used Linear regression.

```
r2 = r2_score(y_actual, y_pred)
mae = mean_absolute_error(y_actual, y_pred)
rmse = np.sqrt(mean_squared_error(y_actual, y_pred))

print(f'R2: {r2:.4f} (higher is better, 0.6-0.8 is decent here)')
print(f'MAE: ${mae:,.0f} (average $ error)')
print(f'RMSE: ${rmse:,.0f}')
```

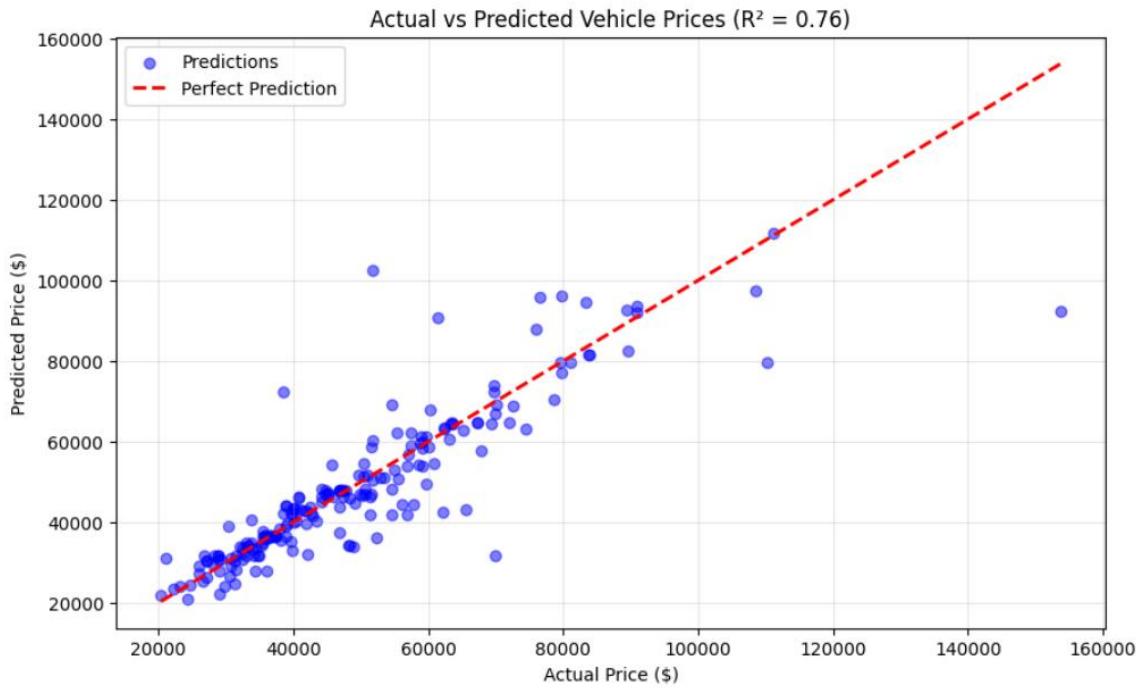
R2: 0.7598 (higher is better, 0.6-0.8 is decent here)
MAE: \$5,004 (average \$ error)
RMSE: \$9,317

- Calculated the R² score and found out its an average score.
- Mean average error is MAE
- Root mean square is RMSE

```
    print("\nSome predictions vs actual (first 10):")
    results = pd.DataFrame({
        'Actual': y_actual.round(0).astype(int),
        'Predicted': y_pred.round(0).astype(int)
    })
    print(results.head(10))
```

| | Actual | Predicted |
|-----|--------|-----------|
| 201 | 83940 | 81492 |
| 556 | 51803 | 60245 |
| 176 | 56105 | 44464 |
| 952 | 37335 | 36416 |
| 66 | 28860 | 31744 |
| 505 | 22260 | 23560 |
| 768 | 29111 | 27936 |
| 561 | 60080 | 58822 |
| 614 | 42150 | 31922 |
| 160 | 45038 | 47290 |

- A comparison between actual data and the data predicted by the model.



- Used Matplotlib to plot the data with the best fit line.
- As you can see , there are anomalies however most of the data is near or on the line.
- However , the R^2 score can be improved and accuracy can be improved.
- Random Forrest Regressor is used.

```
[ ]: from sklearn.ensemble import RandomForestRegressor

model = RandomForestRegressor(n_estimators=100, random_state=42)
model.fit(X_train, y_train)
y_pred_log = model.predict(X_test)
y_pred = np.expm1(y_pred_log)
```

```

r2 = r2_score(y_actual, y_pred)
mae = mean_absolute_error(y_actual, y_pred)
rmse = np.sqrt(mean_squared_error(y_actual, y_pred))

print(f"R²: {r2:.4f} (higher is better, 0.6-0.8 is decent here)")
print(f"MAE: ${mae:,.0f} (average $ error)")
print(f"RMSE: ${rmse:,.0f}")

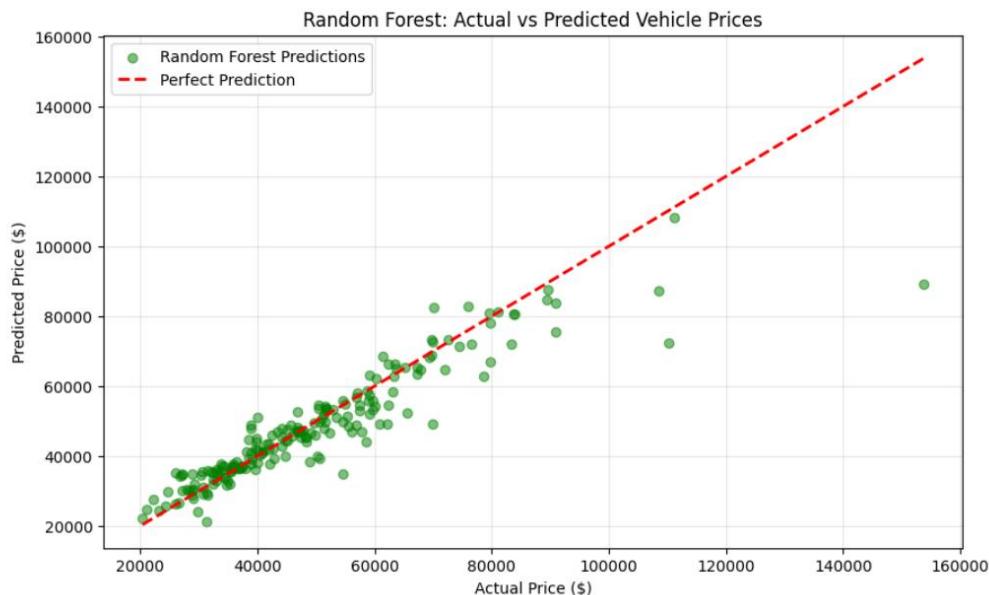
```

R²: 0.8403 (higher is better, 0.6-0.8 is decent here)

MAE: \$4,147 (average \$ error)

RMSE: \$7,597

- Giving an R² score of 0.84 which is 84%.



(3) MOBILE PHONE PRICE PREDICTION

- Similar to the Vehicle Price prediction project almost similar however , I used a different model rather than the Linear Regression , is mainly because to show diversity and the learning through out the unified mentor internship.
- So the model I used here is DECISION TREE MODEL.

```
[10]: df = pd.read_csv('dataset.csv')
[11]: df.head()
[11]:   battery_power  blue  clock_speed  dual_sim  fc  four_g  int_memory  m_dep  mobile_wt  n_cores ...  px_height  px_width  ram  sc_h  sc_w  talk_time  three_g  tot
0          842      0        2.2       0  1      0        7     0.6      188       2 ...        20      756    2549      9      7     19      0
1         1021      1        0.5       1  0      1       53     0.7      136       3 ...       905    1988    2631     17      3      7      1
2          563      1        0.5       1  2      1       41     0.9      145       5 ...      1263    1716    2603     11      2      9      1
3          615      1        2.5       0  0      0       10     0.8      131       6 ...      1216    1786    2769     16      8     11      1
4         1821      1        1.2       0  13     1       44     0.6      141       2 ...      1208    1212    1411      8      2     15      1
```

5 rows × 21 columns

```
[8]: df.isnull().sum()
```

```
[8]: battery_power      0
blue                  0
clock_speed           0
dual_sim              0
fc                    0
four_g                0
int_memory             0
m_dep                 0
mobile_wt              0
n_cores               0
pc                    0
px_height              0
px_width              0
ram                   0
sc_h                  0
sc_w                  0
talk_time              0
three_g                0
touch_screen            0
wifi                  0
price_range             0
dtype: int64
```

- Found out that the dataset had no null values, hence no data cleaning is required like in the previous project.

```
[16]: df['price_range'].value_counts()
```

```
[16]: price_range
      1    500
      2    500
      3    500
      0    500
Name: count, dtype: int64
```

- Main reason for using decision tree here is that the Project is a Multi Class Classification. Which contained 3 categories where the previous one , is Continuous data.

```
[17]: df.corr()['price_range'].sort_values(ascending=False)
```

```
[17]: price_range      1.000000
      ram          0.917046
      battery_power  0.200723
      px_width     0.165818
      px_height    0.148858
      int_memory   0.044435
      sc_w         0.038711
      pc           0.033599
      three_g      0.023611
      sc_h         0.022986
      fc           0.021998
      talk_time    0.021859
      blue         0.020573
      wifi         0.018785
      dual_sim     0.017444
      four_g       0.014772
      n_cores      0.004399
      m_dep        0.000853
      clock_speed   -0.006606
      mobile_wt     -0.030302
      touch_screen  -0.030411
Name: price_range, dtype: float64
```

- This shows the correlation between the output or the Price range with the other features.
- As you can see , the price range is highly dependent on the Ram which is 91.7% of all compared to the other features.

```
[1]: X = df.drop('price_range', axis=1)
y = df['price_range']

print(f"Features shape: {X.shape}")
print(f"Target shape: {y.shape}")

Features shape: (2000, 20)
Target shape: (2000,)

[2]: X_train, X_test, y_train, y_test = train_test_split(
    X, y,
    test_size=0.2,      # 20% for testing
    random_state=42,    # Same split every time
    stratify=y          # Keep same proportion of each class
)
```

- Splitting the data for the Test and the Train where the Test data size is 20% of the data while 80% of the data is for the Training of the data.

```
[3]: print(f"\nTraining set: {X_train.shape[0]} phones ({X_train.shape[0]/len(X)*100:.0f}%)")
print(f"Test set: {X_test.shape[0]} phones ({X_test.shape[0]/len(X)*100:.0f}%)")

Training set: 1600 phones (80%)
Test set: 400 phones (20%)
```

- 1600 phones for Training and 400 phones for Testing.

```
[25]: dt_model = DecisionTreeClassifier(
    max_depth=10,           # Limit depth to prevent overfitting
    min_samples_split=20,    # Need 20 samples to split
    random_state=42
)
```

- Implementing the model Decision Tree.

```
[27]: y_pred_dt = dt_model.predict(X_test)

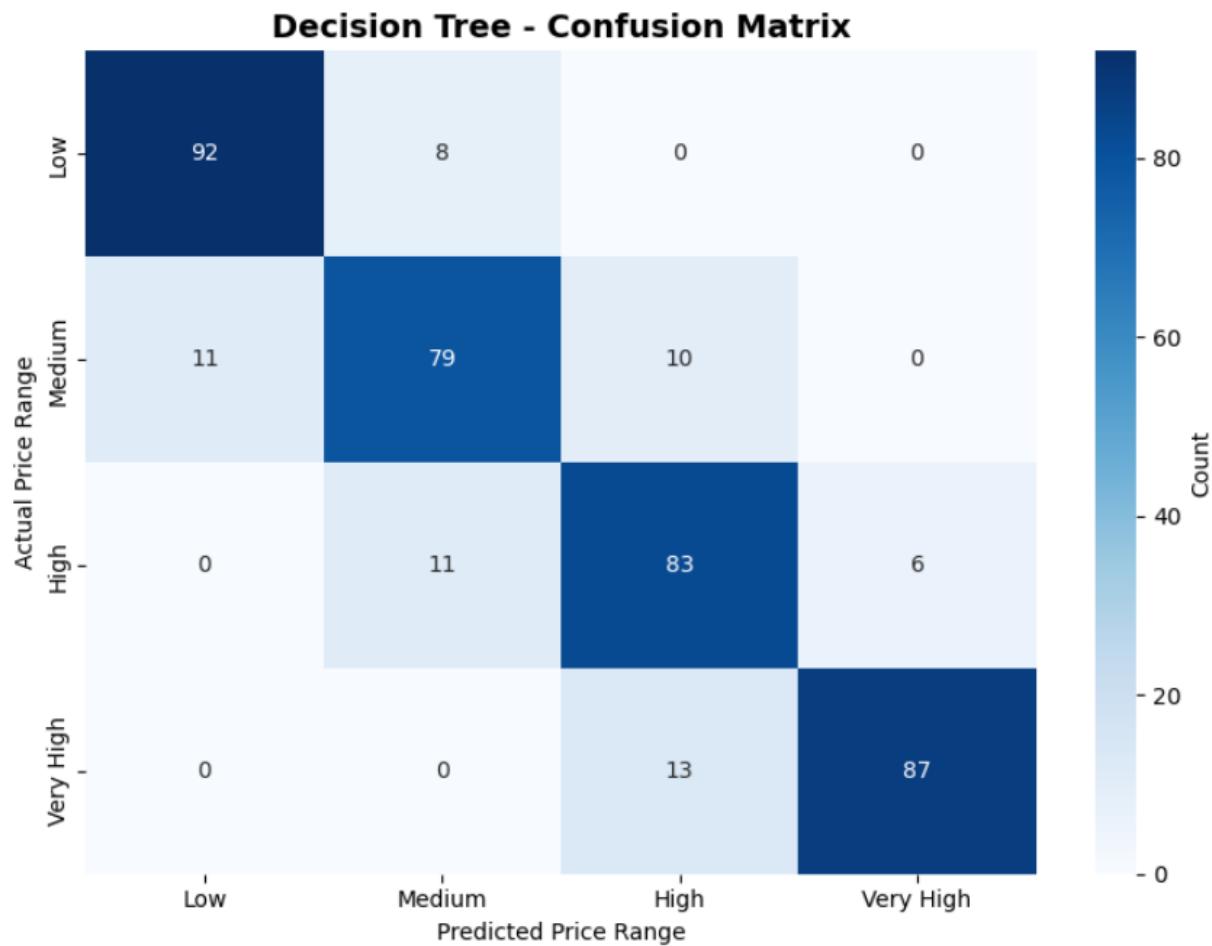
[28]: # Calculate accuracy
      dt_accuracy = accuracy_score(y_test, y_pred_dt)
      print(f"\nDecision Tree Accuracy: {dt_accuracy*100:.2f}%")
```

Decision Tree Accuracy: 85.25%

- Below is the Classification report.

```
=====
CLASSIFICATION REPORT - DECISION TREE
=====
              precision    recall   f1-score   support
Low           0.89      0.92      0.91      100
Medium        0.81      0.79      0.80      100
High          0.78      0.83      0.81      100
Very High     0.94      0.87      0.90      100
accuracy      0.85      0.85      0.85      400
macro avg     0.85      0.85      0.85      400
weighted avg  0.85      0.85      0.85      400
```

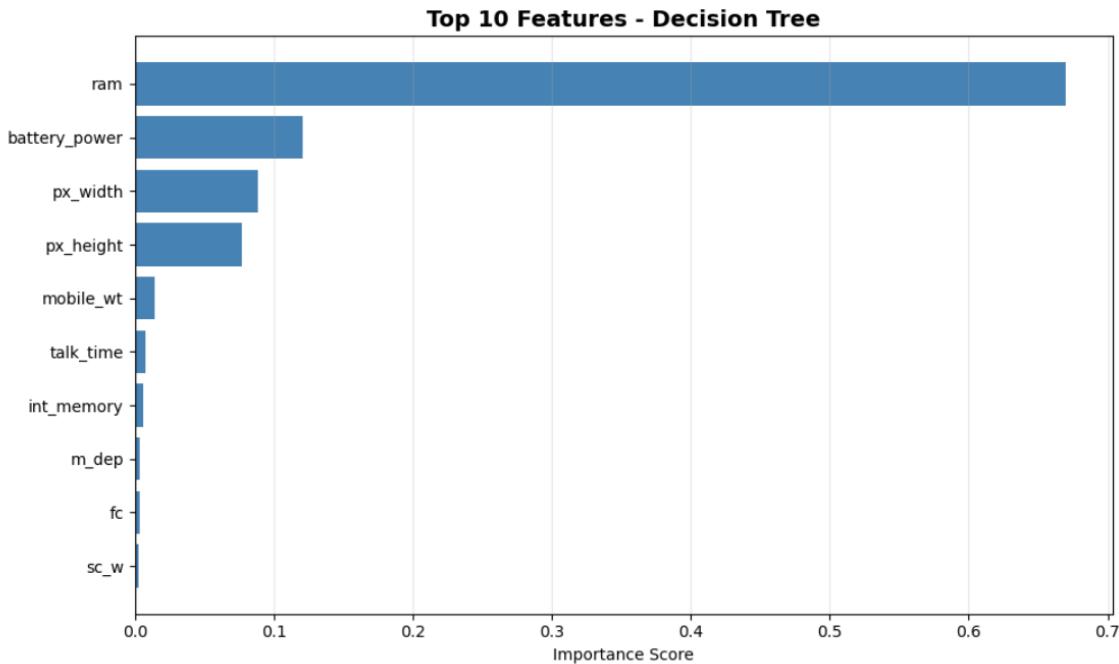
- The confusion Matrix of the above Project is shown below.



FEATURE IMPORTANCE - DECISION TREE

Top 10 Most Important Features:

| | feature | importance |
|----|---------------|------------|
| 13 | ram | 0.669858 |
| 0 | battery_power | 0.120525 |
| 12 | px_width | 0.088476 |
| 11 | px_height | 0.076886 |
| 8 | mobile_wt | 0.013887 |
| 16 | talk_time | 0.007909 |
| 6 | int_memory | 0.005978 |
| 7 | m_dep | 0.003548 |
| 4 | fc | 0.003407 |
| 15 | sc_w | 0.002921 |



- This plot is to show the correlation between the features with the Price range or the importance to the output.
- To make a solid comparison , I implemented Random Forest Classifier model like in the previous project to do the comparison.

```
rf_model = RandomForestClassifier(
    n_estimators=100,           # Build 100 trees
    max_depth=15,              # Each tree can be deeper
    min_samples_split=10,
    random_state=42,
    n_jobs=-1                  # Use all CPU cores
)
```

- Use 100 trees with depth of 15 levels.

```
Random Forest Accuracy: 88.00%
```

```
[38]: # Detailed report
print("\n" + "="*60)
print("CLASSIFICATION REPORT - RANDOM FOREST")
print("="*60)
print(classification_report(y_test, y_pred_rf, target_names=class_names))
```

```
=====
CLASSIFICATION REPORT - RANDOM FOREST
=====
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| Low | 0.94 | 0.96 | 0.95 | 100 |
| Medium | 0.81 | 0.82 | 0.82 | 100 |
| High | 0.82 | 0.79 | 0.81 | 100 |
| Very High | 0.94 | 0.95 | 0.95 | 100 |
| accuracy | | | 0.88 | 400 |
| macro avg | 0.88 | 0.88 | 0.88 | 400 |
| weighted avg | 0.88 | 0.88 | 0.88 | 400 |

- This shows the accuracy 0.88 which is 88%
- Compared with the Decision Tree Model , the Random Forest Classifier has higher accuracy.

DECISION TREE -

```
=====
CLASSIFICATION REPORT - DECISION TREE
=====
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| Low | 0.89 | 0.92 | 0.91 | 100 |
| Medium | 0.81 | 0.79 | 0.80 | 100 |
| High | 0.78 | 0.83 | 0.81 | 100 |
| Very High | 0.94 | 0.87 | 0.90 | 100 |
| accuracy | | | 0.85 | 400 |
| macro avg | 0.85 | 0.85 | 0.85 | 400 |
| weighted avg | 0.85 | 0.85 | 0.85 | 400 |

```
]: correct_dt = (y_pred_dt == y_test).sum()
print(f"\nCorrect predictions: {correct_dt}/{len(y_test)}")
```

```
Correct predictions: 341/400
```

RANDOM FOREST CLASSIFIER

```
=====
```

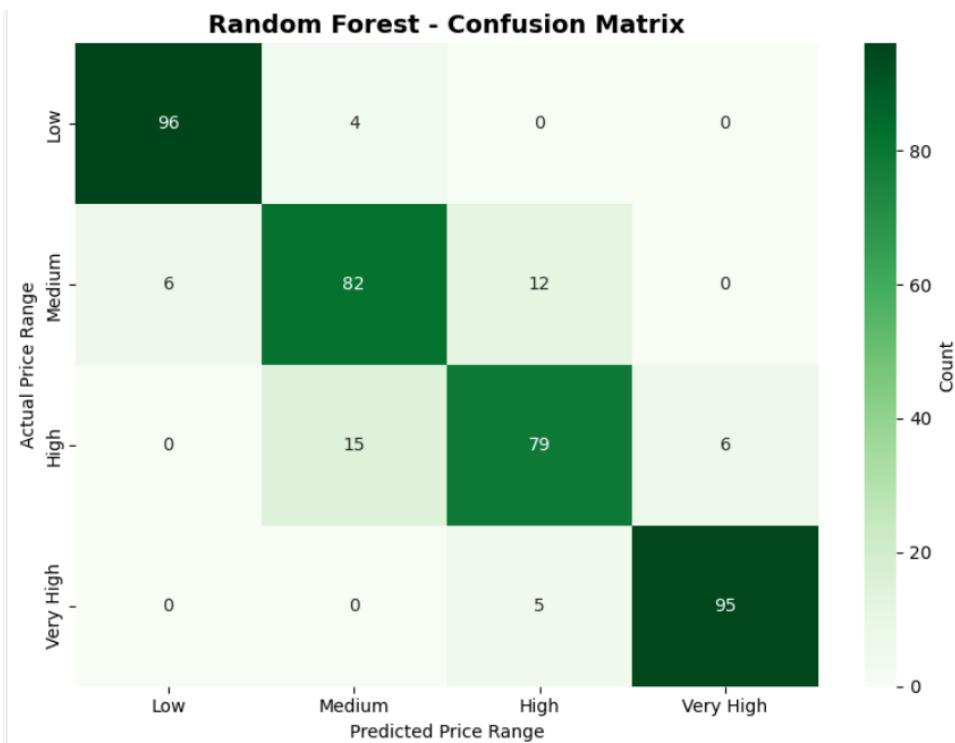
CLASSIFICATION REPORT - RANDOM FOREST

```
=====
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| Low | 0.94 | 0.96 | 0.95 | 100 |
| Medium | 0.81 | 0.82 | 0.82 | 100 |
| High | 0.82 | 0.79 | 0.81 | 100 |
| Very High | 0.94 | 0.95 | 0.95 | 100 |
| accuracy | | | 0.88 | 400 |
| macro avg | 0.88 | 0.88 | 0.88 | 400 |
| weighted avg | 0.88 | 0.88 | 0.88 | 400 |

```
]: # How many correct?  
correct_rf = (y_pred_rf == y_test).sum()  
print(f"\nCorrect predictions: {correct_rf}/{len(y_test)}")
```

Correct predictions: 352/400



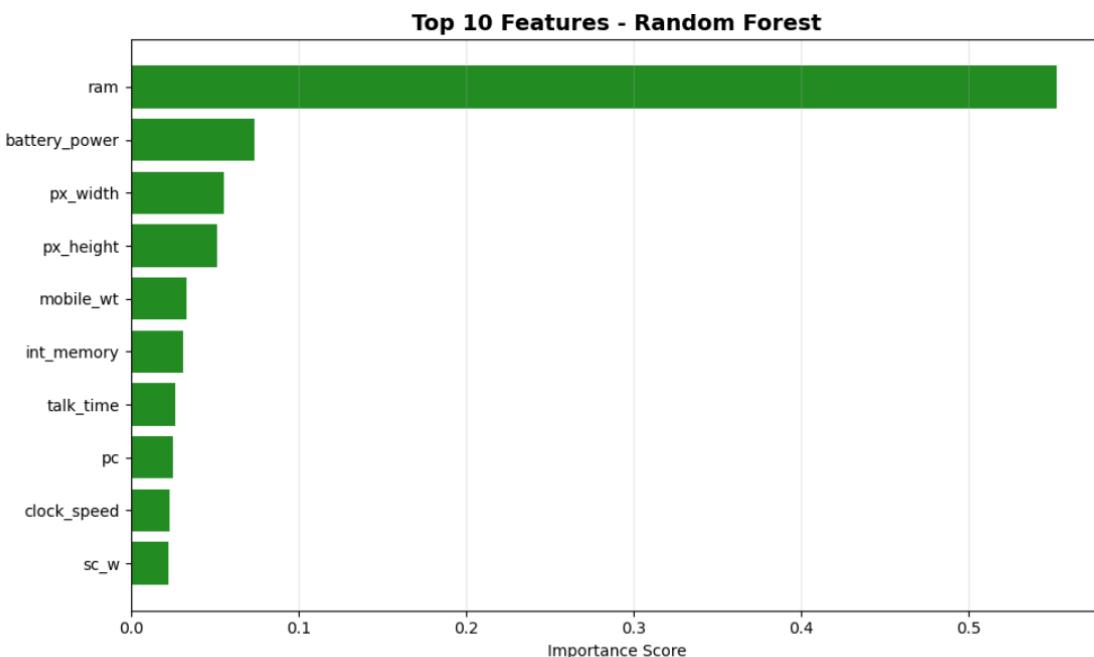
=====

FEATURE IMPORTANCE - RANDOM FOREST

=====

Top 10 Most Important Features:

| | feature | importance |
|----|---------------|------------|
| 13 | ram | 0.552550 |
| 0 | battery_power | 0.074008 |
| 12 | px_width | 0.055601 |
| 11 | px_height | 0.051614 |
| 8 | mobile_wt | 0.033430 |
| 6 | int_memory | 0.031225 |
| 16 | talk_time | 0.026599 |
| 10 | pc | 0.025067 |
| 2 | clock_speed | 0.023410 |
| 15 | sc_w | 0.022805 |

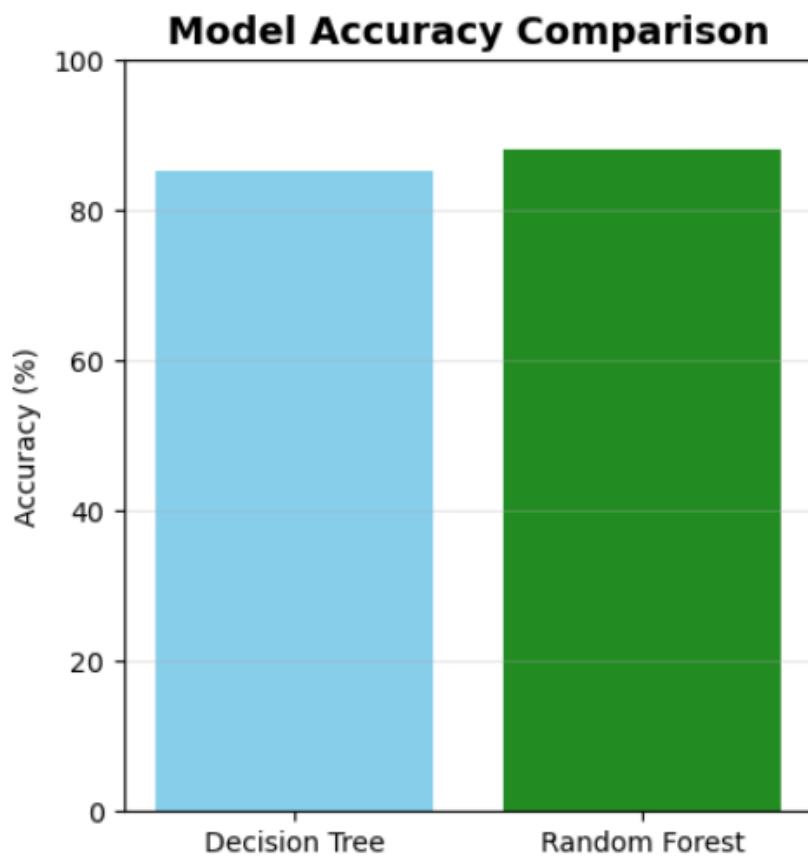


- Like in the previous model , Plot some graphs in order to find the correlation between the price and the features.

COMPARISON BETWEEN MODELS

```
[45]: # Plot accuracy comparison
plt.figure(figsize=(10, 5))

plt.subplot(1, 2, 1)
bars = plt.bar(comparison['Model'], comparison['Accuracy'],
               color=['skyblue', 'forestgreen'])
plt.ylabel('Accuracy (%)')
plt.title('Model Accuracy Comparison', fontsize=14, weight='bold')
plt.ylim(0, 100)
plt.grid(axis='y', alpha=0.3)
```



- Therefore,
 - Random Forrest Classifier has a better accuracy than Decision Tree by 2.75%.

(4) DETECT HEART DISEASE PREDICTION

- Using Logistic Regression Model.
- Mainly because to explore the models learned in Unified Mentor

```
: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import (confusion_matrix, classification_report,
                             accuracy_score, precision_score, recall_score,
                             f1_score, roc_curve, roc_auc_score)

: df = pd.read_csv('dataset.csv')

: df.head()

:   age sex chest pain type resting bp s cholesterol fasting blood sugar resting ecg max heart rate exercise angina oldpeak ST slope target
: 0 40 1 2 140 289 0 0 172 0 0.0 1 0
: 1 49 0 3 160 180 0 0 156 0 1.0 2 1
: 2 37 1 2 130 283 0 1 98 0 0.0 1 0
: 3 48 0 4 138 214 0 0 108 1 1.5 2 1
: 4 54 1 3 150 195 0 0 122 0 0.0 1 0

[3]: df = pd.read_csv('dataset.csv')

[4]: df.head()

[4]:   age sex chest pain type resting bp s cholesterol fasting blood sugar resting ecg max heart rate exercise angina oldpeak ST slope target
: 0 40 1 2 140 289 0 0 172 0 0.0 1 0
: 1 49 0 3 160 180 0 0 156 0 1.0 2 1
: 2 37 1 2 130 283 0 1 98 0 0.0 1 0
: 3 48 0 4 138 214 0 0 108 1 1.5 2 1
: 4 54 1 3 150 195 0 0 122 0 0.0 1 0

[6]: df.isnull().sum()

[6]: age          0
      sex          0
      chest pain type 0
      resting bp s    0
      cholesterol    0
      fasting blood sugar 0
      resting ecg     0
      max heart rate 0
      exercise angina 0
      oldpeak        0
      ST slope       0
      target         0
      dtype: int64
```

- Using Logistic regression model is basically because this is just a Binary Classification. Either 0 or 1.

```

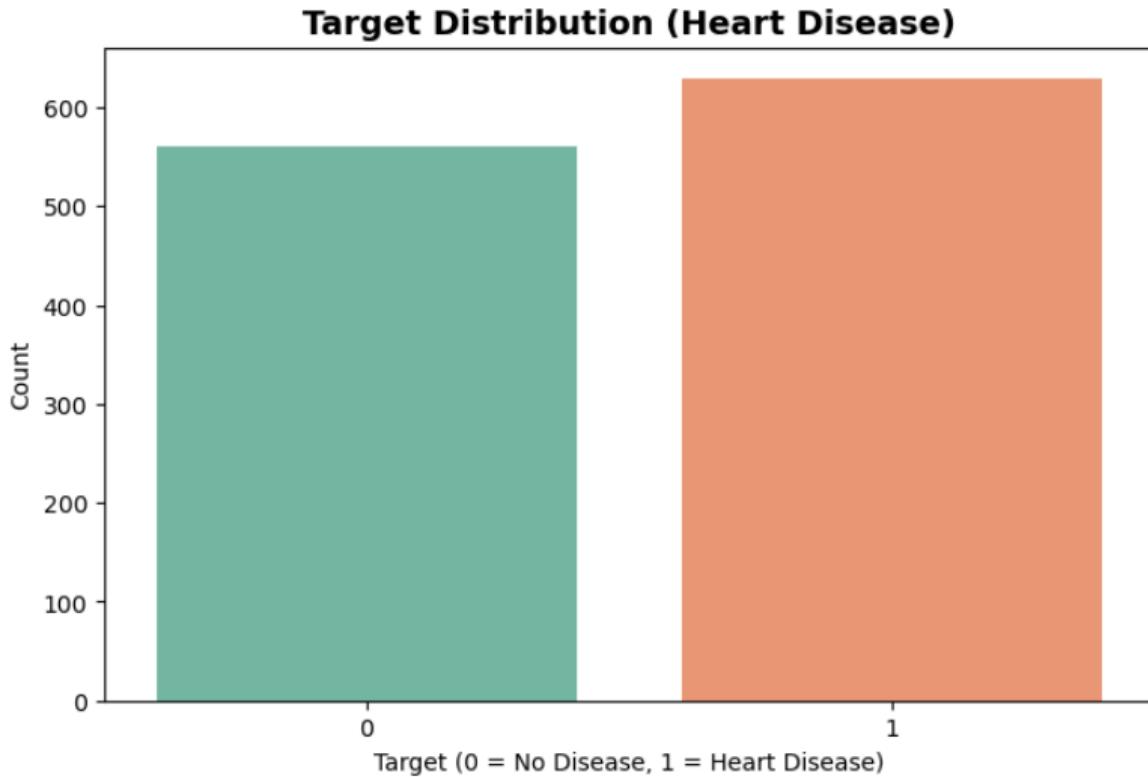
print("Target Distribution:")
print(df['target'].value_counts())
print(f"\nClass 0 (No Disease): {(df['target']==0).sum()} ({((df['target']==0).sum())/len(df)*100:.1f}%)")
print(f"Class 1 (Heart Disease): {(df['target']==1).sum()} ({((df['target']==1).sum())/len(df)*100:.1f}%)")

Target Distribution:
target
1    629
0    561
Name: count, dtype: int64

Class 0 (No Disease): 561 (47.1%)
Class 1 (Heart Disease): 629 (52.9%)

```

- Class 0 or No disease is 47% and Class 1 is 52% which shows that the data set is already balanced hence no modifications needed for the model.
- If the dataset is imbalanced, the model finds it hard to train cause there is no sufficient data for one of the output categories.



```

correlations = df.corr()['target'].sort_values(ascending=False)
print("Features most correlated with Heart Disease:")
print(correlations)

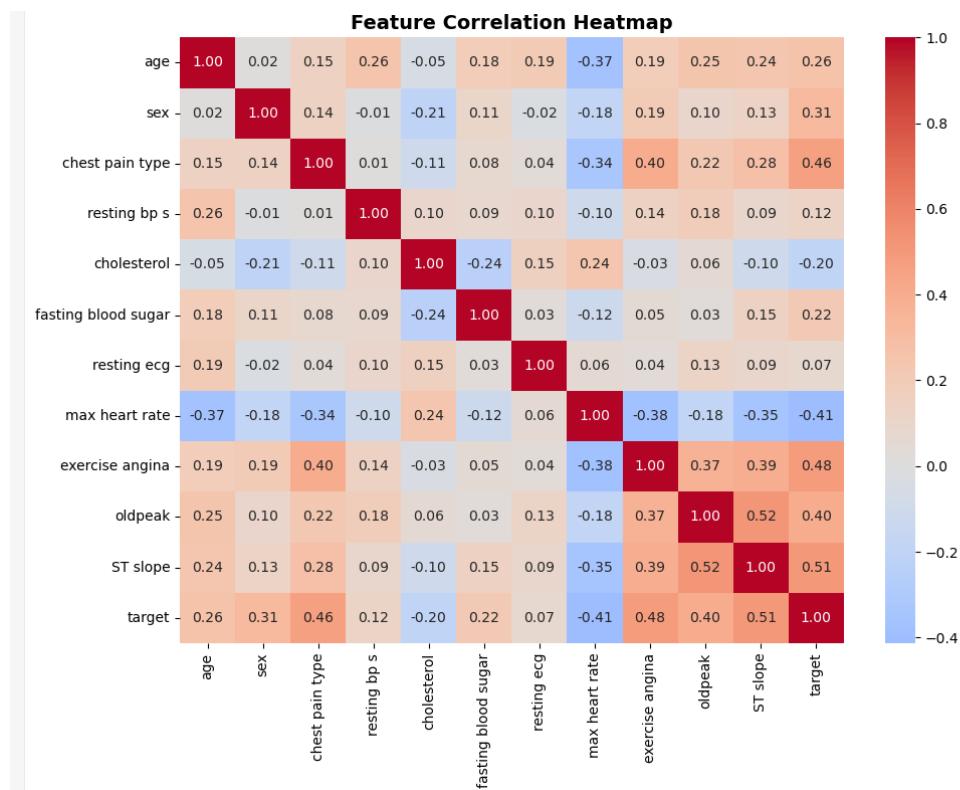
```

Features most correlated with Heart Disease:

| | |
|---------------------|-----------|
| target | 1.000000 |
| ST slope | 0.505608 |
| exercise angina | 0.481467 |
| chest pain type | 0.460127 |
| oldpeak | 0.398385 |
| sex | 0.311267 |
| age | 0.262029 |
| fasting blood sugar | 0.216695 |
| resting bp s | 0.121415 |
| resting ecg | 0.073059 |
| cholesterol | -0.198366 |
| max heart rate | -0.413278 |

Name: target, dtype: float64

- A comparison or shows the relationship/correlation between the target and the other features.



- For more comparison , I implemented a correlation heatmap which shows the correlations between the target and other features. The negative high values shows inversely impact while positive high values shows direct impact and less values shows almost no impact.

```
[13]: # Cell: Feature Scaling
print("🔧 Scaling features using StandardScaler...")
print("\n💡 Why scaling is important:")
print("  • ST slope ranges: 1-3")
print("  • Max heart rate ranges: 71-202")
print("  • Different scales slow down Gradient Descent!")
print("  • StandardScaler makes all features mean=0, std=1")

scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
X_scaled = pd.DataFrame(X_scaled, columns=feature_columns)

print("\nScaling complete!")
print("\nBefore vs After (first 3 samples):")
print("\nBefore:")
print(X.head(3))
print("\nAfter (standardized):")
print(X_scaled.head(3))
```

- Using standard scaler so that values comes under 0-1 rather than big values making the model easy to learn and train.

Before vs After (first 3 samples):

| | age | sex | chest pain | type | resting bp | s | cholesterol | fasting blood sugar | \ |
|---|-----|-----|------------|------|------------|---|-------------|---------------------|---|
| 0 | 40 | 1 | | 2 | 140 | | 289 | | 0 |
| 1 | 49 | 0 | | 3 | 160 | | 180 | | 0 |
| 2 | 37 | 1 | | 2 | 130 | | 283 | | 0 |

| | resting ecg | max heart rate | exercise | angina | oldpeak | ST slope |
|---|-------------|----------------|----------|--------|---------|----------|
| 0 | 0 | 172 | | 0 | 0.0 | 1 |
| 1 | 0 | 156 | | 0 | 1.0 | 2 |
| 2 | 1 | 98 | | 0 | 0.0 | 1 |

| | age | sex | chest pain | type | resting bp | s | cholesterol | \ |
|---|-----------|-----------|------------|-----------|------------|---|-------------|---|
| 0 | -1.466728 | 0.555995 | | -1.318351 | 0.427328 | | 0.775674 | |
| 1 | -0.504600 | -1.798576 | | -0.248932 | 1.516587 | | -0.299512 | |
| 2 | -1.787437 | 0.555995 | | -1.318351 | -0.117301 | | 0.716489 | |

| | fasting blood sugar | resting ecg | max heart rate | exercise | angina | \ |
|---|---------------------|-------------|----------------|----------|-----------|---|
| 0 | -0.520929 | -0.802672 | 1.265039 | | -0.795219 | |
| 1 | -0.520929 | -0.802672 | 0.637758 | | -0.795219 | |
| 2 | -0.520929 | 0.346762 | -1.636136 | | -0.795219 | |

```
[14]: # Cell: Split data
X_train, X_test, y_train, y_test = train_test_split(
    X_scaled, y,
    test_size=0.2,
    random_state=42,
    stratify=y # Keeps 47-53 balance in both sets!
)

print("Data split complete!")
print(f"\n Training set: {X_train.shape[0]} samples")
print(f" Class 0: {sum(y_train==0)} ({(sum(y_train==0)/len(y_train)*100:.1f}%)")
print(f" Class 1: {sum(y_train==1)} ({(sum(y_train==1)/len(y_train)*100:.1f}%)")

print(f"\n Testing set: {X_test.shape[0]} samples")
print(f" Class 0: {sum(y_test==0)} ({(sum(y_test==0)/len(y_test)*100:.1f}%)")
print(f" Class 1: {sum(y_test==1)} ({(sum(y_test==1)/len(y_test)*100:.1f}%)")
```

Data split complete!

```
Training set: 952 samples
Class 0: 449 (47.2%)
Class 1: 503 (52.8%)

Testing set: 238 samples
Class 0: 112 (47.1%)
Class 1: 126 (52.9%)
```

- Splitting the data between the training and the test datasets for plotting.

```
[15]: model = LogisticRegression(max_iter=1000, random_state=42)
model.fit(X_train, y_train)
```

[15]: ▾ LogisticRegression ⓘ ⓘ

► Parameters

- Logistic regression model.

```

print("Model Parameters:")
print(f"Intercept: {model.intercept_[0]:.4f}")

print("\nFeature Coefficients:")
for feature, coef in zip(feature_columns, model.coef_[0]):
    print(f"{feature:25s}: {coef:.4f}")

```

Model Parameters:
Intercept: 0.2344

Feature Coefficients:

| | | |
|---------------------|---|---------|
| age | : | 0.2162 |
| sex | : | 0.6148 |
| chest pain type | : | 0.6910 |
| resting bp s | : | 0.1296 |
| cholesterol | : | -0.2561 |
| fasting blood sugar | : | 0.3888 |
| resting ecg | : | 0.0072 |
| max heart rate | : | -0.2212 |
| exercise angina | : | 0.5083 |
| oldpeak | : | 0.3996 |
| ST slope | : | 0.7798 |

- This shows the y intercept which is Theta0 and then other features like X1θ1 , X2 θ2 etc

```

print("Sample Predictions:")
for i in range(10):
    print(f"Actual: {y_test.iloc[i]}, Predicted: {y_test_pred[i]}, Probability: {y_test_prob[i]:.2f}")

```

Sample Predictions:

| |
|--|
| Actual: 1, Predicted: 1, Probability: 0.96 |
| Actual: 0, Predicted: 0, Probability: 0.45 |
| Actual: 1, Predicted: 1, Probability: 0.94 |
| Actual: 0, Predicted: 1, Probability: 0.53 |
| Actual: 0, Predicted: 0, Probability: 0.09 |
| Actual: 1, Predicted: 1, Probability: 0.89 |
| Actual: 0, Predicted: 0, Probability: 0.02 |
| Actual: 0, Predicted: 0, Probability: 0.22 |
| Actual: 0, Predicted: 0, Probability: 0.01 |
| Actual: 1, Predicted: 1, Probability: 0.90 |

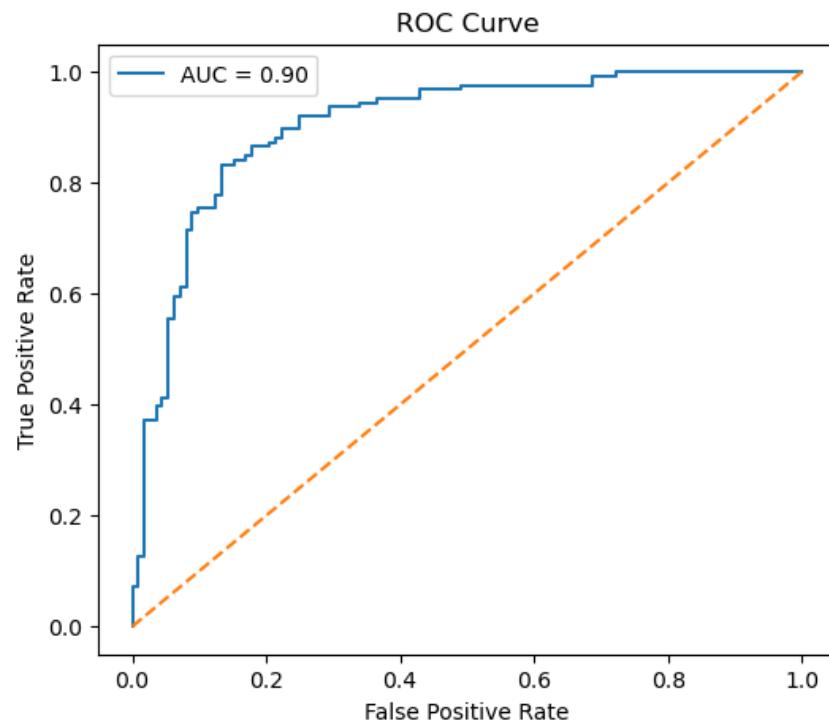
- Shows the predictions between the actual value and the predicted by the model.

- Predictions like 0.53 gives the model more error because slight difference leads to different prediction where actual is 0 but predicted as 1 since $0.5 > 0.53$

```
from sklearn.metrics import roc_curve, auc

fpr, tpr, _ = roc_curve(y_test, y_test_prob)
roc_auc = auc(fpr, tpr)

plt.figure(figsize=(6, 5))
plt.plot(fpr, tpr, label=f"AUC = {roc_auc:.2f}")
plt.plot([0, 1], [0, 1], linestyle="--")
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("ROC Curve")
plt.legend()
plt.show()
```



- Implemented ROC Curve to show more better between the True positive and the False positives rates.

False Positive Rate = False Positives / (False Positives + True Negatives)

➔ Lower the better

True Positive Rate = True Positives / (True Positives + False Negatives)

➔ Higher the better.

- ❖ Blue Line is the Model, shows how well the model distinguish between the healthy and sick or between 1 and 0.
- ❖ Orange Dashed Line is the Random Classifier, this is model guess randomly. Any model which is above this line is better.
- ❖ AUC = Area under the curve which is Total area under BLUE LINE